

PyTorch TTFR (ms, log scale)

- Lazy (demand-only)
- Lazy + fixed prefetch
- Lazy + adaptive
- ★ Real two-machine (LAN)

crossover
100-150 μ s

real LAN, 311 μ s

*congestion-bound:
prefetch harmful*

*latency-bound:
prefetch wins*

Injected round-trip time (μ s, log scale; leftmost = loopback)

