

PyTorch TTFR (ms, log scale)

- Lazy (demand-only)
- Lazy + fixed prefetch
- Lazy + adaptive
- ★ Real two-machine (LAN)

crossover
~125 μ s RTT

latency-bound
(prefetch wins)
real LAN
311 μ s

congestion-bound
(prefetch harmful)

Injected round-trip time (μ s, log scale; leftmost = loopback)

